

MSRA SH Triton Study Group

2. Triton Softmax Kernel

2025/07/11

What is Triton

- A language
 - An open-source Python-like programming language
 - Enables researchers with no CUDA experience to write highly efficient GPU code
- A compiler
 - Compile Triton language code to executable CUDA kernels
 - With higher readability than other existing DSLs.

Triton language and programming model

Triton

- Kernel
- Program
- INVISIBLE
- INVISIBLE

CUDA

- Kernel
- Thread Block
- (32 threads)
- Thread

Hardware

- GPU
- Streaming Multiprocessor
- Warp
- Lane

Triton language and programming model


- Vectorized and parallel programming
 - Define an array in registers or shared memory across a thread block
`tl.zeros()`, `tl.arange()`, etc.
 - The compiler automatically choose the best parallelization scheme in a program (thread block)
- Torch-like functions for tensor operation
 - Element-wise basic computation: `+`, `-`, `*`, `/`, `tl.exp()`, `tl.log()`, etc.
 - Reduce on an axis: `tl.sum()`, `tl.max()`, etc.
 - Matrix multiplication: `tl.dot()`

Triton Softmax Kernel

- Equivalent to better softmax kernel in the previous class
- 1 warp per row
 - 1 program (thread blocks) = 4 warps = 128 threads
- For each row i :
 - Load all X_{ij} in once
 - Reduce to find M_i
 - Reduce to calc S_i
 - Store Y_{ij}

Triton Softmax Kernel

TB#0, T#0	TB#0, T#1	...	TB#0, T#31
TB#0, T#32	TB#0, T#33	...	TB#0, T#63
TB#0, T#64	TB#0, T#65	...	TB#0, T#95
TB#0, T#96	TB#0, T#97	...	TB#0, T#127
TB#1, T#0	TB#1, T#1	...	TB#1, T#31
TB#1, T#32	TB#1, T#33	...	TB#1, T#63
TB#1, T#64	TB#1, T#65	...	TB#1, T#95
TB#1, T#96	TB#1, T#97	...	TB#1, T#127
TB#2, T#0	TB#2, T#1	...	TB#2, T#31
...



```
template<int VALS_PER_THREAD> // Each thread (lane) processes VALS_PER_THREAD values
__global__ void better_softmax_kernel(float* x, float* y, int batch_size) {
    // Current warp index in the thread block
    const int warp_idx = threadIdx.x / WARP_SIZE;
    // Current thread lane index within the warp
    const int lane_idx = threadIdx.x % WARP_SIZE;
```



```
# Triton kernel for softmax operation
@triton.jit
def softmax_kernel(
    x_ptr, # Pointer to input tensor
    y_ptr, # Pointer to output tensor
    batch_size,
    hidden_dim,
    BLOCK_SIZE_M: tl.constexpr, # Block size for M (batch_size) dimension
    BLOCK_SIZE_N: tl.constexpr, # Block size for N (hidden_dim) dimension
):
    # Equivalent to blockIdx in CUDA
    pid = tl.program_id(0)
```

```
// Each warp processes one row
const int row_idx = blockIdx.x * num_warps + warp_idx;
// Boundary check
if (row_idx >= batch_size) return;

// Offset for contiguous memory access in a warp
const int offset = row_idx * (WARP_SIZE * VALS_PER_THREAD) + lane_idx * MEM_ACCESS_WIDTH;
```



```
# Boundary check
start_m = pid * BLOCK_SIZE_M
if start_m >= batch_size:
    return

# Offsets for M and N dimensions
offs_m = start_m + tl.arange(0, BLOCK_SIZE_M)
offs_n = tl.arange(0, BLOCK_SIZE_N)
```



```

// Allocate VALS_PER_THREAD floats in registers
float tmp_val[VALS_PER_THREAD];

// Load VALS_PER_THREAD values from global memory into registers
#pragma unroll
for (int i = 0; i < VALS_PER_THREAD; i += MEM_ACCESS_WIDTH) {
    // Vectorized memory access using float4
    reinterpret_cast<float4*>(&tmp_val[i])[0] =
        reinterpret_cast<float4*>(&x[offset + i * WARP_SIZE])[0];
}

```



```

# Mask invalid M offsets when BLOCK_SIZE_M > 1
mask = offs_m[:, None] < batch_size

# Load input tensor
x = tl.load(x_ptr + offs_m[:, None] * hidden_dim + offs_n[None, :], mask=mask)

```

```


// Find the maximum value in the thread's values
float max_val = -FLT_MAX;
#pragma unroll
for (int i = 0; i < VALS_PER_THREAD; i++) {
    max_val = max(max_val, tmp_val[i]);
}

// Reduce the maximum value across all threads in the warp
#pragma unroll
for (int laneMask = 1; laneMask < WARP_SIZE; laneMask <= 1) {
    max_val = max(max_val, __shfl_xor_sync(WARP_MASK, max_val, laneMask));
}

// Calculate exponential values and sum of exponentials in the thread's values
float sum_exp = 0.0f;
#pragma unroll
for (int i = 0; i < VALS_PER_THREAD; i++) {
    tmp_val[i] = expf(tmp_val[i] - max_val);
    sum_exp += tmp_val[i];
}

// Reduce the sum of exponentials across all threads in the warp
#pragma unroll
for (int laneMask = 1; laneMask < WARP_SIZE; laneMask <= 1) {
    sum_exp += __shfl_xor_sync(WARP_MASK, sum_exp, laneMask);
}

```



```

# Compute softmax
m = tl.max(x, axis=1, keep_dims=True)
e = tl.exp(x - m)
s = tl.sum(e, axis=1, keep_dims=True)
y = e / s

```

```
// Write VALS_PER_THREAD values registers to global memory
#pragma unroll
for (int i = 0; i < VALS_PER_THREAD; i += MEM_ACCESS_WIDTH) {
    // Vectorized memory access using float4
    reinterpret_cast<float4*>(&y[offset + i * WARP_SIZE])[0] =
        reinterpret_cast<float4*>(&tmp_val[i])[0];
}
```



```
# Store to the output tensor
tl.store(y_ptr + offs_m[:, None] * hidden_dim + offs_n[None, :], y, mask=mask)
```

```
// Number of warps in a thread block
const int num_warps = 4;
// Grid size (= number of thread blocks)
int num_blocks = (batch_size + num_warps - 1) / num_warps;

const dim3 dimBlock(num_warps * WARP_SIZE);
const dim3 dimGrid(num_blocks);
```



```
# Program (thread block) size = 4 warps = 128 threads
num_warps = 4

# Each program (thread block) processes 4 rows
BLOCK_SIZE_M = 4
BLOCK_SIZE_N = hidden_dim
assert hidden_dim & (hidden_dim - 1) == 0, "Hidden dimension must be a power of 2"
```

```

// Launch the kernel with the appropriate vals_per_thread
if (hidden_dim % WARP_SIZE != 0) {
    throw std::runtime_error("Unsupported hidden dimension size.");
}
int vals_per_thread = hidden_dim / WARP_SIZE;
if (vals_per_thread == 4) {
    better_softmax_kernel<4><<<dimGrid, dimBlock>>>(X.data_ptr<float>(), Y.data_ptr<float>(), batch_size);
} else if (vals_per_thread == 8) {
    better_softmax_kernel<8><<<dimGrid, dimBlock>>>(X.data_ptr<float>(), Y.data_ptr<float>(), batch_size);
} else if (vals_per_thread == 16) {
    better_softmax_kernel<16><<<dimGrid, dimBlock>>>(X.data_ptr<float>(), Y.data_ptr<float>(), batch_size);
} else if (vals_per_thread == 32) {
    better_softmax_kernel<32><<<dimGrid, dimBlock>>>(X.data_ptr<float>(), Y.data_ptr<float>(), batch_size);
} else if (vals_per_thread == 64) {
    better_softmax_kernel<64><<<dimGrid, dimBlock>>>(X.data_ptr<float>(), Y.data_ptr<float>(), batch_size);
} else if (vals_per_thread == 128) {
    better_softmax_kernel<128><<<dimGrid, dimBlock>>>(X.data_ptr<float>(), Y.data_ptr<float>(), batch_size);
} else {
    throw std::runtime_error("Unsupported hidden dimension size.");
}

```



```

# Launch the Triton kernel
grid = (triton.cdiv(batch_size, BLOCK_SIZE_M), )
softmax_kernel[grid](
    x, y, batch_size, hidden_dim,
    BLOCK_SIZE_M=BLOCK_SIZE_M, BLOCK_SIZE_N=BLOCK_SIZE_N,
    num_warps=num_warps,
)

```

Triton Softmax Kernel

- Almost the same speed with CUDA better_softmax_kernel()

```
Batch size: 8765, Hidden dim: 4096  
Torch Softmax Latency: 0.202 ms  
CUDA Softmax Latency: 0.186 ms  
Triton Softmax Latency: 0.184 ms
```

- Try different tile with low cost

```
# Each program (thread block) processes 1 rows  
BLOCK_SIZE_M = 1
```

```
Batch size: 8765, Hidden dim: 4096  
Torch Softmax Latency: 0.202 ms  
CUDA Softmax Latency: 0.186 ms  
Triton Softmax Latency: 0.182 ms (BLOCK_SIZE_M=1)
```

Homework

- Play with [TritonStudyGroup/2_Triton_Softmax at main · Starmys/TritonStudyGroup](#)
- Questions:
 - Implement a Triton softmax kernel for arbitrary hidden dimension
 - Which things have changed when we set BLOCK_SIZE_M=1?

Reading Materials

- Triton introduction and tutorials
 - [Triton-lang/triton: Development repository for the Triton language and compiler](#)
 - [Introducing Triton: Open-source GPU programming for neural networks | OpenAI](#)
 - [Tutorials — Triton documentation](#)
- Other AI kernel compilers
 - [apache/tvm: Open deep learning compiler stack for cpu, gpu and specialized accelerators](#)
 - [ThunderKittens: Simple, Fast, and Adorable AI Kernels](#)
 - [TileLang: A Composable Tiled Programming Model for AI Systems](#)