


MSRA SH Triton Study Group

6. Triton Flash-Attention

2025/09/05

Flash-Attention 2 Forward

- Tile \mathbf{Q}, \mathbf{O} into $\frac{N_Q}{T_Q}$ blocks $\mathbf{Q}_i, \mathbf{O}_i$
- Tile \mathbf{K}, \mathbf{V} into $\frac{N_K}{T_K}$ blocks $\mathbf{K}_j, \mathbf{V}_j$
- Parallel for $1 \leq i \leq \frac{N_Q}{T_Q}$
 - in Registers: $\mathbf{O}_i \leftarrow (0)_{\text{float32}}^{T_Q \times D}, M_i \leftarrow (-\infty)_{\text{float32}}^{T_Q}, L_i \leftarrow (0)_{\text{float32}}^{T_Q}$
 - For $1 \leq j \leq \frac{N_K}{T_K}$
 - $\mathbf{S}_{ij} \leftarrow \frac{\mathbf{Q}_i \mathbf{K}_j^T}{\sqrt{D}}, M_i^{\text{local}} \leftarrow \max_{-1}(\mathbf{S}_{ij})$
 - $M_i^{\text{new}} \leftarrow \max(M_i, M_i^{\text{local}}), \mathbf{P}_{ij} \leftarrow \frac{\exp(\mathbf{S}_{ij} - M_i^{\text{new}})}{L_i}, L_i^{\text{local}} \leftarrow \text{sum}_{-1}(\exp(\mathbf{S}_{ij} - M_i^{\text{new}}))$
 - $\alpha \leftarrow \exp(M_i - M_i^{\text{new}}), L_i \leftarrow \alpha L_i + L_i^{\text{local}}, M_i \leftarrow M_i^{\text{new}}$
 - $\mathbf{O}_i \leftarrow \alpha \mathbf{O}_i + \mathbf{P}_{ij} \mathbf{V}_j$
 - Save $\mathbf{O}_i \leftarrow \frac{\mathbf{O}_i}{L_i}$ (to float16) and $LSE_i \leftarrow M_i + \log(L_i)$ to GPU HBM



```

batch_size, num_tokens, num_heads, head_dim = q.shape
batch_size_k, num_k_tokens, num_k_heads, head_dim_k = k.shape
batch_size_v, num_v_tokens, num_v_heads, head_dim_v = v.shape
assert batch_size == batch_size_k and batch_size_k == batch_size_v
assert num_tokens == num_k_tokens and num_k_tokens == num_v_tokens
assert num_heads == num_k_heads and num_k_heads == num_v_heads
assert head_dim == head_dim_k and head_dim_k == head_dim_v
assert head_dim in {16, 32, 64, 128}

sm_scale = sm_scale or head_dim ** (-0.5)
o = torch.empty_like(q)
lse = torch.empty((batch_size, num_heads, num_tokens), dtype=torch.float32, device=q.device)

grid = lambda META: (triton.cdiv(num_tokens, META['BLOCK_SIZE_M']), num_heads * batch_size, 1)
flash_attn_fwd_kernel[grid](
    q, k, v, o, lse,
    q.stride(0), q.stride(2), q.stride(1), q.stride(3),
    k.stride(0), k.stride(2), k.stride(1), k.stride(3),
    v.stride(0), v.stride(2), v.stride(1), v.stride(3),
    o.stride(0), o.stride(2), o.stride(1), o.stride(3),
    lse.stride(0), lse.stride(1), lse.stride(2),
    sm_scale, num_tokens, num_heads,
    CAUSAL=causal, BLOCK_SIZE_D=head_dim,
)

```

```
@triton.jit
```

```
def flash_attn_fwd_kernel(
```

```
    q_ptr, k_ptr, v_ptr, o_ptr, lse_ptr,  
    stride_qb, stride_qh, stride_qm, stride_qd,  
    stride_kb, stride_kh, stride_kn, stride_kd,  
    stride_vb, stride_vh, stride_vn, stride_vd,  
    stride_ob, stride_oh, stride_om, stride_od,  
    stride_lb, stride_lh, stride_lm,  
    sm_scale, num_tokens, num_heads, CAUSAL: tl.constexpr,  
    BLOCK_SIZE_M: tl.constexpr,  
    BLOCK_SIZE_N: tl.constexpr,  
    BLOCK_SIZE_D: tl.constexpr,
```

```
):
```

```
    start_m = tl.program_id(0) * BLOCK_SIZE_M  
    off_h = tl.program_id(1) % num_heads  
    off_b = tl.program_id(1) // num_heads
```

```
    # Initialize offsets
```

```
    offs_m = start_m + tl.arange(0, BLOCK_SIZE_M)  
    offs_n = tl.arange(0, BLOCK_SIZE_N)  
    offs_d = tl.arange(0, BLOCK_SIZE_D)  
    mask_m = offs_m < num_tokens
```

```
    q_ptrs = q_ptr + off_b * stride_qb + off_h * stride_qh + offs_m[:, None] * stride_qm + offs_d[None, :] * stride_qd  
    k_ptrs = k_ptr + off_b * stride_kb + off_h * stride_kh + offs_n[:, None] * stride_kn + offs_d[None, :] * stride_kd  
    v_ptrs = v_ptr + off_b * stride_vb + off_h * stride_vh + offs_n[:, None] * stride_vn + offs_d[None, :] * stride_vd  
    o_ptrs = o_ptr + off_b * stride_ob + off_h * stride_oh + offs_m[:, None] * stride_om + offs_d[None, :] * stride_od  
    lse_ptrs = lse_ptr + off_b * stride_lb + off_h * stride_lh + offs_m * stride_lm
```

```

m = tl.zeros([BLOCK_SIZE_M], dtype=tl.float32) - float("inf")
l = tl.zeros([BLOCK_SIZE_M], dtype=tl.float32)
o = tl.zeros([BLOCK_SIZE_M, BLOCK_SIZE_D], dtype=tl.float32)
q = tl.load(q_ptrs, mask=mask_m[:, None], other=0.0)

# Scale sm_scale by log_2(e) and use 2^x instead of exp
q = (q * sm_scale * 1.4426950408889634).to(q_ptr.type.element_ty)

# Split the main loop into 2 parts
if CAUSAL:
    start, mid, end = 0, start_m // BLOCK_SIZE_N * BLOCK_SIZE_N, tl.minimum(start_m + BLOCK_SIZE_M, num_tokens)
else:
    start, mid, end = 0, num_tokens // BLOCK_SIZE_N * BLOCK_SIZE_N, num_tokens

# Main loop part 1: no mask at first
o, m, l = _attn_fwd_loop(
    q, k_ptrs, v_ptrs, o, m, l,
    offs_m, offs_n, stride_kn, stride_vn,
    start, mid, BLOCK_SIZE_N=BLOCK_SIZE_N,
    CAUSAL=False, MASK_N=False,
)

# Main loop part 2: causal mask and kv data mask for last blocks
o, m, l = _attn_fwd_loop(
    q, k_ptrs, v_ptrs, o, m, l,
    offs_m, offs_n, stride_kn, stride_vn,
    mid, end, BLOCK_SIZE_N=BLOCK_SIZE_N,
    CAUSAL=CAUSAL, MASK_N=True,
)

```

```
@triton.jit
def _attn_fwd_loop(
    q, k_ptrs, v_ptrs, o, m, l,
    offs_m, offs_n, stride_kn, stride_vn,
    start, end, BLOCK_SIZE_N: tl.constexpr,
    CAUSAL: tl.constexpr, MASK_N: tl.constexpr,
):
    for start_n in range(start, end, BLOCK_SIZE_N):

        # Load K, V
        ...

        # Calc S <- Q @ K^T, mask S if causal
        ...

        # Calc new row max M' <- max(M, max(s)), P <- exp(S - M'), local sum exp L1 <- sum(P)
        ...

        # Update L <- L * exp(M - M') + L1, M <- M'
        ...

        # Update O <- O * exp(M - M') + P @ V
        ...

    return o, m, l
```

```

# Load K, V
if MASK_N:
    mask_n = start_n + offs_n < end
    k = tl.load(k_ptrs + start_n * stride_kn, mask=mask_n[:, None], other=0.0)
    v = tl.load(v_ptrs + start_n * stride_vn, mask=mask_n[:, None], other=0.0)
else:
    k = tl.load(k_ptrs + start_n * stride_kn)
    v = tl.load(v_ptrs + start_n * stride_vn)

# Calc S <- Q @ K^T, mask S if causal
s = tl.dot(q, k.T)
if CAUSAL:
    s = tl.where(offs_m[:, None] >= start_n + offs_n[None, :], s, float("-inf"))
elif MASK_N:
    s = tl.where(mask_n[None, :], s, float("-inf"))

# Calc new row max M' <- max(M, max(s)), P <- exp(S - M'), local sum exp L1 <- sum(P)
m_new = tl.maximum(m, tl.max(s, 1))
p = tl.math.exp2(s - m_new[:, None])
l_local = tl.sum(p, 1)


# Update L <- L * exp(M - M') + L1, M <- M'
alpha = tl.math.exp2(m - m_new)
l = l * alpha + l_local
m = m_new

# Update O <- O * exp(M - M') + P @ V
o = o * alpha[:, None] + tl.dot(p.to(v.type.element_ty), v)

```

Flash-Attention 2 Backward

- Preprocess: $\Delta_i = \sum_{k=0}^D d\mathbf{O}_{ik} \odot \mathbf{O}_{ik}$, $d\mathbf{Q} \leftarrow (0)_{\text{float32}}^{N_Q \times D}$
- Parallel for $1 \leq j \leq \frac{N_K}{T_K}$
 - in Registers: $d\mathbf{K}_j \leftarrow (0)_{\text{float32}}^{T_K \times D}$, $d\mathbf{V}_j \leftarrow (0)_{\text{float32}}^{T_K \times D}$
 - For $1 \leq i \leq \frac{N_Q}{T_Q}$
 - $\mathbf{P}_{ij} \leftarrow \exp\left(\frac{\mathbf{Q}_i \mathbf{K}_j^T}{\sqrt{D}} - LSE_i\right)$, $d\mathbf{V}_j \leftarrow d\mathbf{V}_j + \mathbf{P}_{ij}^T d\mathbf{O}_i$
 - $d\mathbf{S}_{ij} \leftarrow \mathbf{P}_{ij} \odot (d\mathbf{O}_i \mathbf{V}_j^T - \Delta_i)$, $d\mathbf{K}_j \leftarrow d\mathbf{K}_j + \frac{d\mathbf{S}_{ij}^T \mathbf{Q}_i}{\sqrt{D}}$
 - $d\mathbf{Q}_i \leftarrow d\mathbf{Q}_i + \frac{d\mathbf{S}_{ij} \mathbf{K}_j}{\sqrt{D}}$ by atomic_add()
 - Save $d\mathbf{K}_j$ and $d\mathbf{V}_j$ to GPU HBM (to float16)
- Convert $d\mathbf{Q}$ to float16



Preprocess: $\Delta_i = \sum_{k=0}^D d\mathbf{o}_{ik} \odot \mathbf{o}_{ik}$

```
@torch.compile
def flash_attn_bwd_calc_delta(
    o: torch.Tensor,  # [batch_size, num_tokens, num_heads, head_dim]
    do: torch.Tensor, # [batch_size, num_tokens, num_heads, head_dim]
):
    return torch.sum((o * do).swapaxes(1, 2), dim=-1, dtype=torch.float32)
```

```
class TritonFlashAttention(torch.autograd.Function):
```

```
    @staticmethod
```

```
    def forward(
```

```
        ctx: torch.autograd.function.FunctionCtx,
```

```
        q: torch.Tensor,
```

```
        k: torch.Tensor,
```

```
        v: torch.Tensor,
```

```
        sm_scale: float,
```

```
        causal: bool,
```

```
    ):
```

```
        batch_size, num_tokens, num_heads, head_dim = q.shape
```

```
        # Check shapes
```

```
        ...
```

```
        # Launch forward kernel
```

```
        o, lse = ...
```

```
        ctx.batch_size = batch_size
```

```
        ctx.num_heads = num_heads
```

```
        ctx.num_tokens = num_tokens
```

```
        ctx.head_dim = head_dim
```

```
        ctx.sm_scale = sm_scale
```

```
        ctx.causal = causal
```

```
        ctx.save_for_backward(q, k, v, o, lse)
```

```
        return o
```

```
    @staticmethod
```

```
    def backward(ctx, do: torch.Tensor):
```

```
        q, k, v, o, lse = ctx.saved_tensors
```

```
        delta = flash_attn_bwd_calc_delta(o, do)
```

```
        dq = torch.zeros_like(q, dtype=torch.float32)
```

```
        dk = torch.empty_like(k)
```

```
        dv = torch.empty_like(v)
```

```
        grid = lambda META: (
```

```
            triton.cdiv(ctx.num_tokens, META['BLOCK_SIZE_N']),
```

```
            ctx.num_heads * ctx.batch_size,
```

```
        )
```

```
        flash_attn_bwd_kernel[grid](
```

```
            q, k, v, lse, delta, dq, dk, dv, do,
```

```
            q.stride(0), q.stride(2), q.stride(1), q.stride(3),
```

```
            k.stride(0), k.stride(2), k.stride(1), k.stride(3),
```

```
            v.stride(0), v.stride(2), v.stride(1), v.stride(3),
```

```
            o.stride(0), o.stride(2), o.stride(1), o.stride(3),
```

```
            lse.stride(0), lse.stride(1), lse.stride(2),
```

```
            delta.stride(0), delta.stride(1), delta.stride(2),
```

```
            ctx.sm_scale, ctx.num_tokens, ctx.num_heads,
```

```
            CAUSAL=ctx.causal, BLOCK_SIZE_D=ctx.head_dim,
```

```
            SKIP_DQ=ctx.split_bwd,
```

```
        )
```

```
        return dq.to(q.dtype), dk, dv, None, None, None
```

@triton.jit

def flash_attn_bwd_kernel(

q_ptr, k_ptr, v_ptr, lse_ptr, delta_ptr,

dq_ptr, dk_ptr, dv_ptr, do_ptr,

stride_qb, stride_qh, stride_qm, stride_qd,

stride_kb, stride_kh, stride_kn, stride_kd,

stride_vb, stride_vh, stride_vn, stride_vd,

stride_ob, stride_oh, stride_om, stride_od,

stride_lb, stride_lh, stride_lm,

stride_db, stride_dh, stride_dm,

sm_scale, num_tokens, num_heads, CAUSAL: tl.constexpr,

BLOCK_SIZE_M: tl.constexpr, BLOCK_SIZE_N: tl.constexpr, BLOCK_SIZE_D: tl.constexpr,

):

start_n = tl.program_id(0) * BLOCK_SIZE_N

off_h = tl.program_id(1) % num_heads

off_b = tl.program_id(1) // num_heads

Initialize offsets

offs_m = tl.arange(0, BLOCK_SIZE_M)

offs_n = start_n + tl.arange(0, BLOCK_SIZE_N)

offs_d = tl.arange(0, BLOCK_SIZE_D)

mask_n = offs_n < num_tokens

q_ptrs = q_ptr + off_b * stride_qb + off_h * stride_qh + offs_m[:, None] * stride_qm + offs_d[None, :] * stride_qd

k_ptrs = k_ptr + off_b * stride_kb + off_h * stride_kh + offs_n[:, None] * stride_kn + offs_d[None, :] * stride_kd

v_ptrs = v_ptr + off_b * stride_vb + off_h * stride_vh + offs_n[:, None] * stride_vn + offs_d[None, :] * stride_vd

dq_ptrs = dq_ptr + off_b * stride_qb + off_h * stride_qh + offs_m[None, :] * stride_qm + offs_d[:, None] * stride_qd

dk_ptrs = dk_ptr + off_b * stride_kb + off_h * stride_kh + offs_n[:, None] * stride_kn + offs_d[None, :] * stride_kd

dv_ptrs = dv_ptr + off_b * stride_vb + off_h * stride_vh + offs_n[:, None] * stride_vn + offs_d[None, :] * stride_vd

do_ptrs = do_ptr + off_b * stride_ob + off_h * stride_oh + offs_m[:, None] * stride_om + offs_d[None, :] * stride_od

lse_ptrs = lse_ptr + off_b * stride_lb + off_h * stride_lh + offs_m * stride_lm

delta_ptrs = delta_ptr + off_b * stride_db + off_h * stride_dh + offs_m * stride_dm

```

dk = tl.zeros([BLOCK_SIZE_N, BLOCK_SIZE_D], dtype=tl.float32)
dv = tl.zeros([BLOCK_SIZE_N, BLOCK_SIZE_D], dtype=tl.float32)
k = (tl.load(k_ptrs, mask=mask_n[:, None], other=0.0) * (sm_scale * 1.4426950408889634)).to(k_ptr.type.element_ty)
v = tl.load(v_ptrs, mask=mask_n[:, None], other=0.0)

# Split the main loop into 3 parts
if CAUSAL:
    causal_start = start_n // BLOCK_SIZE_M * BLOCK_SIZE_M
    causal_end = tl.minimum(causal_start + tl.maximum(BLOCK_SIZE_M, BLOCK_SIZE_N), num_tokens)
    full_start = causal_end
    full_end = num_tokens // BLOCK_SIZE_M * BLOCK_SIZE_M
    mask_start = tl.maximum(full_end, causal_end)
    mask_end = num_tokens
else:
    causal_start, causal_end = 0, 0
    full_start, full_end = 0, num_tokens // BLOCK_SIZE_M * BLOCK_SIZE_M
    mask_start, mask_end = full_end, num_tokens

# Main loop part 3: mask last rows that exceed the sequence length
dk, dv = _attn_bwd_loop(mask_start, mask_end, CAUSAL=False, MASK_M=True)
# Main loop part 2: no mask
dk, dv = _attn_bwd_loop(full_start, full_end, CAUSAL=False, MASK_M=False)
# Main loop part 1: causal mask
dk, dv = _attn_bwd_loop(causal_start, causal_end, CAUSAL=True, MASK_M=True)

# Write back dK and dV
tl.store(dk_ptrs, (dk * sm_scale).to(dk_ptr.type.element_ty), mask=mask_n[:, None])
tl.store(dv_ptrs, dv.to(dv_ptr.type.element_ty), mask=mask_n[:, None])

```

```

@triton.jit
def _attn_bwd_loop(
    k, v, dk, dv, dq_ptrs, q_ptrs, do_ptrs, lse_ptrs, delta_ptrs,
    offs_m, offs_n, stride_qm, stride_om, stride_lm, stride_dm,
    start, end, BLOCK_SIZE_M: tl.constexpr, CAUSAL: tl.constexpr, MASK_M: tl.constexpr,
):
    # Reverse loop to maximize L2 cache hit rate when causal
    for start_m in range(tl.cdiv(end, BLOCK_SIZE_M) * BLOCK_SIZE_M - BLOCK_SIZE_M, start - BLOCK_SIZE_M, -BLOCK_SIZE_M):

        # Load Q, dO, LSE (= M + log(L)), Δ (= sum(dO * O))
        ...

        # Calc P <- exp(Q @ K^T - M) / L
        ...

        # Update dV <- dV + P^T @ dO
        ...

        # Calc dS <- P * (dO @ V^T - Δ)
        ...

        # Update dQ <- dQ + dS @ K by atomic_add
        ...

        # Update dK <- dK + dS^T @ Q
        ...

    return dk, dv

```

```

# Load Q, dO, LSE (= M + log(L)), Δ (= sum(dO * O))
q = tl.load(q_ptrs + start_m * stride_qm)
do = tl.load(do_ptrs + start_m * stride_om)
lse = tl.load(lse_ptrs + start_m * stride_lm)
delta = tl.load(delta_ptrs + start_m * stride_dm)

# Calc P <- exp(Q @ K^T - M) / L
pT = tl.math.exp2(tl.dot(k, q.T) - lse[None, :])
if CAUSAL:
    causal_mask = start_m + offs_m[None, :] >= offs_n[:, None]
    pT = tl.where(causal_mask, pT, 0.0)

# Update dV <- dV + P^T @ dO
dv += tl.dot(pT.to(do.type.element_ty), do)

# Calc dS <- P * (dO @ V^T - Δ)
dsT = (pT * (tl.dot(v, do.T) - delta[None, :])).to(q.type.element_ty)


# Update dQ <- dQ + dS @ K by atomic_add
dqT = tl.dot(k.T, dsT.to(k.type.element_ty)) * 0.6931471824645996
tl.atomic_add(dq_ptrs + start_m * stride_om, dqT, sem='relaxed')

# Update dK <- dK + dS^T @ Q
dk += tl.dot(dsT.to(q.type.element_ty), q)

```

Calculate dQ separately using a new kernel

- Possible reasons:
 - Triton atomic_add() not good
 - Triton compile is not optimized
- Parallel for $1 \leq i \leq \frac{N_Q}{T_Q}$
 - in Registers: $dQ_i \leftarrow (0)_{\text{float32}}^{T_Q \times D}$
 - For $1 \leq j \leq \frac{N_K}{T_K}$
 - $P_{ij} \leftarrow \exp\left(\frac{Q_i K_j^T}{\sqrt{D}} - LSE_i\right)$
 - $dS_{ij} \leftarrow P_{ij} \odot (dO_i V_j^T - \Delta_i)$
 - $dQ_i \leftarrow dQ_i + \frac{dS_{ij} K_j}{\sqrt{D}}$
 - Save dQ_i to GPU HBM



Latency (Triton 3.2.0)

- Forward
 - Much faster than torch
 - 10%~15% slower than official
- Backward
 - atomic_add: ~80% slower
 - split_bwd: ~40% slower

```
=====
Problem size: (1, 4096, 32, 128); Causal = False
          fwd      bwd
triton_atomic_add    1.653951    7.951271
triton_split_bwd     1.661846    5.531596
flash_official       1.450124    4.171456
torch_naive          12.335835   19.094136
-----
Problem size: (1, 4096, 32, 128); Causal = True
          fwd      bwd
triton_atomic_add    0.968440    4.335084
triton_split_bwd     0.968935    3.191221
flash_official       0.838765    2.319285
torch_naive          16.054613   24.706221
-----
Problem size: (1, 4321, 32, 128); Causal = False
          fwd      bwd
triton_atomic_add    1.947054    8.969808
triton_split_bwd     1.945600    6.765008
flash_official       1.681893    4.926344
torch_naive          21.459967   31.271936
-----
Problem size: (1, 4321, 32, 128); Causal = True
          fwd      bwd
triton_atomic_add    1.082380    4.805748
triton_split_bwd     1.092376    3.625090
flash_official       0.946758    2.642283
torch_naive          26.221909   37.962410
=====
```


Homework

- Play with [TritonStudyGroup/6_Triton_Flash_Attn at main · Starmys/TritonStudyGroup](#)
- Support Grouped-Query Attention
- Implement Triton block-sparse Flash-Attention with block size of 128
- Output the attention scores **after pooling**
- Why do we split the bwd main loop into 3 parts (but not 2 parts)?
- Profile Triton flash-attention forward and explain why it is slow
- Calculate the FLOPs of `triton_split_bwd`

Reading Materials

- Triton flash attention example: [Fused Attention — Triton documentation](#)