

# MSRA SH Triton Study Group

## 7. Triton Fused Kernels

2025/09/12

# Problem #1: Fused SiLU

- Definition:

$$\text{silu}(x, y) = \frac{xy}{1 + \exp(x)}$$

- Full element-wise function

```
Problem size: (4096, 4096)
      naive_silu  fused_silu  compiled_silu
latency    0.701396    0.129117    0.129992
```

```

@triton.jit
def _fused_silu_kernel(
    x_ptr, y_ptr, M, N,
    stride_xm, stride_xn,
    stride_ym, stride_yn,
    BLOCK_SIZE_M: tl.constexpr, BLOCK_SIZE_N: tl.constexpr,
):
    pid_m = tl.program_id(axis=1)
    pid_n = tl.program_id(axis=0)
    offs_m = pid_m * BLOCK_SIZE_M + tl.arange(0, BLOCK_SIZE_M)
    offs_n = pid_n * BLOCK_SIZE_N + tl.arange(0, BLOCK_SIZE_N)
    mask = offs_m < M
    offs_m %= M
    x1_ptrs = x_ptr + offs_m[:, None] * stride_xm + offs_n[None, :] * stride_xn
    x3_ptrs = x1_ptrs + N * stride_xn
    y_ptrs = y_ptr + offs_m[:, None] * stride_ym + offs_n[None, :] * stride_yn
    x1 = tl.load(x1_ptrs).to(tl.float32)
    x3 = tl.load(x3_ptrs).to(tl.float32)
    y = x1 * (1 / (1 + tl.exp(-x1))) * x3
    tl.store(y_ptrs, y.to(y_ptr.type.element_ty), mask=mask[:, None])

```

# Problem #2: Fused MoE Routing

- Definition

```
def moe_routing(
    hidden_states: torch.Tensor,
    gating_weight: torch.Tensor,
    num_experts: int,
    top_k: int,
):
    num_tokens, hidden_dim = hidden_states.shape
    router_logits = torch.nn.functional.linear(hidden_states, gating_weight, bias=None)
    router_logits = torch.nn.functional.softmax(router_logits, dim=1, dtype=torch.float32)
    routing_weights, selected_experts = torch.topk(router_logits, top_k, dim=-1)
    routing_weights = routing_weights.to(hidden_states.dtype)
    expert_mask = torch.nn.functional.one_hot(selected_experts.flatten(), num_classes=num_experts)
    expert_cnt = expert_mask.sum(dim=0)
    expert_idx = expert_mask.argsort(dim=0, descending=True)[:num_tokens].T.contiguous()
    return router_logits, routing_weights, expert_cnt, expert_idx, selected_experts
```

```

@triton.jit
def _moe_routing_kernel(
    a_ptr, # [M, K]
    b_ptr, # [E, K]
    c_ptr, # [M, E]
    w_ptr, # [M, T]
    t_ptr, # [M, T]
    cnt_ptr, # [E] in [0, M)
    idx_ptr, # [E, M] in [0, T * M)
    M, E, K, T,
    stride_am, stride_ak,
    stride_be, stride_bk,
    stride_cm, stride_ce,
    stride_wm, stride_we,
    stride_tm, stride_te,
    BLOCK_SIZE_M: tl.constexpr, BLOCK_SIZE_K: tl.constexpr,
    BLOCK_SIZE_E: tl.constexpr, BLOCK_SIZE_T: tl.constexpr,
):
    pid_m = tl.program_id(axis=0) # pid_e = 0

    if pid_m * BLOCK_SIZE_M >= M:
        return

    offs_m = (pid_m * BLOCK_SIZE_M + tl.arange(0, BLOCK_SIZE_M)) % M
    offs_e = tl.arange(0, BLOCK_SIZE_E)
    offs_k = tl.arange(0, BLOCK_SIZE_K)

    a_ptrs = a_ptr + offs_m[:, None] * stride_am + offs_k[None, :] * stride_ak
    b_ptrs = b_ptr + offs_e[None, :] * stride_be + offs_k[:, None] * stride_bk

    # GeMM
    accumulator = tl.zeros([BLOCK_SIZE_M, BLOCK_SIZE_E], dtype=tl.float32)
    for k in range(0, K, BLOCK_SIZE_K):
        a = tl.load(a_ptrs)
        b = tl.load(b_ptrs)
        accumulator += tl.dot(a, b)
        a_ptrs += BLOCK_SIZE_K * stride_ak
        b_ptrs += BLOCK_SIZE_K * stride_bk

```

```

# Softmax
c_max = tl.max(accumulator, axis=1)
c_exp = tl.math.exp(accumulator - c_max[:, None])
c_sum = tl.sum(c_exp, axis=1)
c = c_exp / c_sum[:, None]

# Save Routing Weights
offs_m = pid_m * BLOCK_SIZE_M + tl.arange(0, BLOCK_SIZE_M)
c_ptrs = c_ptr + offs_m[:, None] * stride_cm + offs_e[None, :] * stride_ce
tl.store(c_ptrs, c.to(c_ptr.type.element_ty), mask=mask_m[:, None])

# Top-K
w_ptrs = w_ptr + offs_m * stride_wm
t_ptrs = t_ptr + offs_m * stride_tm
offs_t = tl.arange(0, BLOCK_SIZE_T)
e_idx = tl.zeros([BLOCK_SIZE_M, BLOCK_SIZE_T], dtype=tl.int32)
for k in tl.static_range(BLOCK_SIZE_T):
    max_c, max_idx = tl.max(c, axis=1, return_indices=True)
    c = tl.where(offs_e[None, :] == max_idx[:, None], 0.0, c)
    e_idx = tl.where(offs_t[None, :] == k, max_idx[:, None], e_idx)
    tl.store(w_ptrs, max_c.to(w_ptr.type.element_ty), mask=mask_m)
    tl.store(t_ptrs, max_idx.to(t_ptr.type.element_ty), mask=mask_m)
    w_ptrs += stride_we
    t_ptrs += stride_te

# Histogram
e_idx = tl.reshape(e_idx, [BLOCK_SIZE_M * BLOCK_SIZE_T, ])
mask = tl.reshape(
    tl.broadcast_to(offs_m[:, None] < M, [BLOCK_SIZE_M, BLOCK_SIZE_T]),
    [BLOCK_SIZE_M * BLOCK_SIZE_T, ])
)
m_idx = tl.atomic_add(
    cnt_ptr + e_idx, tl.zeros_like(e_idx) + 1,
    mask=mask, sem='relaxed',
)
)
token_idx = tl.reshape(
    offs_m[:, None] * T + offs_t[None, :],
    [BLOCK_SIZE_M * BLOCK_SIZE_T, ],
)
)
tl.store(idx_ptr + e_idx * M + m_idx, token_idx, mask=mask)

```

# Problem #2: Fused MoE Routing

- Complex workload
  - Matmul
  - Row reduce (softmax, top-K)
  - Custom reduce (atomic add)
- Latency


```
Problem size: (4096, 4096)
      naive_routing  fused_routing  compiled_routing
latency      2.095337      0.097784      2.028817
```

# Problem #3: Sparse Attention

- Definition: add custom mask to attention kernels

# Flash Attention

- Parallel for  $1 \leq i \leq \frac{N_Q}{T_Q}$ 
  - Initialize  $\mathbf{O}_i, M_i, L_i$
  - Load  $\mathbf{Q}_i$  from HBM
  - For  $1 \leq j \leq \frac{N_K}{T_K}$ 
    - Load  $\mathbf{K}_j$  and  $\mathbf{V}_j$  from HBM
    - $\mathbf{S}_{ij} \leftarrow \frac{\mathbf{Q}_i \mathbf{K}_j^T}{\sqrt{D}}, M_i^{\text{local}} \leftarrow \max_{-1}(\mathbf{S}_{ij})$
    - $M_i^{\text{new}} \leftarrow \max(M_i, M_i^{\text{local}}), \mathbf{P}_{ij} \leftarrow \frac{\exp(\mathbf{S}_{ij} - M_i^{\text{new}})}{L_i}, L_i^{\text{local}} \leftarrow \text{sum}_{-1}(\exp(\mathbf{S}_{ij} - M_i^{\text{new}}))$
    - $\alpha \leftarrow \exp(M_i - M_i^{\text{new}}), L_i \leftarrow \alpha L_i + L_i^{\text{local}}, M_i \leftarrow M_i^{\text{new}}$
    - $\mathbf{O}_i \leftarrow \alpha \mathbf{O}_i + \mathbf{P}_{ij} \mathbf{V}_j$
  - Save  $\mathbf{O}_i \leftarrow \frac{\mathbf{O}_i}{L_i}$  and  $LSE_i \leftarrow M_i + \log(L_i)$  to HBM





```

@triton.jit
def _attn_fwd_loop(
    q, k_ptrs, v_ptrs, o, m, l,
    offs_m, offs_n, stride_kn, stride_vn,
    start, end, BLOCK_SIZE_N: tl.constexpr,
):
    for start_n in range(start, end, BLOCK_SIZE_N):

        # Load K, V
        k = tl.load(k_ptrs + start_n * stride_kn)
        v = tl.load(v_ptrs + start_n * stride_vn)

        # Calc S <- Q @ K^T
        ...

        # Calc new row max M' <- max(M, max(s)), P <- exp(S - M'), local sum exp L1 <- sum(P)
        ...

        # Update L <- L * exp(M - M') + L1, M <- M'
        ...


        # Update O <- O * exp(M - M') + P @ V
        ...

    return o, m, l

```

# Block-Sparse Attention

- Parallel for  $1 \leq i \leq \frac{N_Q}{T_Q}$ 
  - Initialize  $\mathbf{O}_i, M_i, L_i$
  - Load  $\mathbf{Q}_i$  from HBM
  - For  $1 \leq j \leq CNT_i$ 
    - Load  $\mathbf{K}_{IDX_{ij}}$  and  $\mathbf{V}_{IDX_{ij}}$  from HBM
    - $\mathbf{S}_{ij} \leftarrow \frac{\mathbf{Q}_i \mathbf{K}_j^T}{\sqrt{D}}, M_i^{\text{local}} \leftarrow \max_{-1}(\mathbf{S}_{ij})$
    - $M_i^{\text{new}} \leftarrow \max(M_i, M_i^{\text{local}}), \mathbf{P}_{ij} \leftarrow \frac{\exp(\mathbf{S}_{ij} - M_i^{\text{new}})}{L_i}, L_i^{\text{local}} \leftarrow \text{sum}_{-1}(\exp(\mathbf{S}_{ij} - M_i^{\text{new}}))$
    - $\alpha \leftarrow \exp(M_i - M_i^{\text{new}}), L_i \leftarrow \alpha L_i + L_i^{\text{local}}, M_i \leftarrow M_i^{\text{new}}$
    - $\mathbf{O}_i \leftarrow \alpha \mathbf{O}_i + \mathbf{P}_{ij} \mathbf{V}_j$
  - Save  $\mathbf{O}_i \leftarrow \frac{\mathbf{O}_i}{L_i}$  and  $LSE_i \leftarrow M_i + \log(L_i)$  to HBM



```

def _attn_fwd_loop(
    q, k_ptrs, v_ptrs, o, m, l,
    offs_m, offs_n, stride_kn, stride_vn,
    block_cnt, block_idx, BLOCK_SIZE_N: tl.constexpr,
):
    for j in range(0, block_cnt):

        # Load K, V
        start_n = tl.load(block_idx + j) * BLOCK_SIZE_N
        k = tl.load(k_ptrs + start_n * stride_kn)
        v = tl.load(v_ptrs + start_n * stride_vn)

        # Calc  $S \leftarrow Q @ K^T$ 
        ...

        # Calc new row max  $M' \leftarrow \max(M, \max(s))$ ,  $P \leftarrow \exp(S - M')$ , local sum exp  $L1 \leftarrow \sum(P)$ 
        ...

        # Update  $L \leftarrow L * \exp(M - M') + L1$ ,  $M \leftarrow M'$ 
        ...


        # Update  $O \leftarrow O * \exp(M - M') + P @ V$ 
        ...

    return o, m, l

```

# PIT Attention

- Parallel for  $1 \leq i \leq \frac{N_Q}{T_Q}$ 
  - Initialize  $\mathbf{O}_i, M_i, L_i$
  - Load  $\mathbf{Q}_i$  from HBM
  - For  $1 \leq j \leq \frac{CNT_i}{T_K}$ 
    - Load  $\mathbf{K}_{IDX_{i,jT_K:(j+1)T_K}}$  and  $\mathbf{V}_{IDX_{i,jT_K:(j+1)T_K}}$  from HBM
    - $\mathbf{S}_{ij} \leftarrow \frac{\mathbf{Q}_i \mathbf{K}_j^T}{\sqrt{D}}, M_i^{\text{local}} \leftarrow \max_{-1}(\mathbf{S}_{ij})$
    - $M_i^{\text{new}} \leftarrow \max(M_i, M_i^{\text{local}}), \mathbf{P}_{ij} \leftarrow \frac{\exp(\mathbf{S}_{ij} - M_i^{\text{new}})}{L_i}, L_i^{\text{local}} \leftarrow \text{sum}_{-1}(\exp(\mathbf{S}_{ij} - M_i^{\text{new}}))$
    - $\alpha \leftarrow \exp(M_i - M_i^{\text{new}}), L_i \leftarrow \alpha L_i + L_i^{\text{local}}, M_i \leftarrow M_i^{\text{new}}$
    - $\mathbf{O}_i \leftarrow \alpha \mathbf{O}_i + \mathbf{P}_{ij} \mathbf{V}_j$
  - Save  $\mathbf{O}_i \leftarrow \frac{\mathbf{O}_i}{L_i}$  and  $LSE_i \leftarrow M_i + \log(L_i)$  to HBM



```

def _attn_fwd_loop(
    q, k_ptrs, v_ptrs, o, m, l,
    offs_m, offs_n, stride_kn, stride_vn,
    col_cnt, col_idx, BLOCK_SIZE_N: tl.constexpr,
):
    for start_n in range(0, col_cnt, BLOCK_SIZE_N):

        # Load K, V
        idx = tl.load(col_idx + start_n + offs_n)
        k = tl.load(k_ptrs + idx * stride_kn)
        v = tl.load(v_ptrs + idx * stride_vn)

        # Calc  $S \leftarrow Q @ K^T$ 
        ...

        # Calc new row max  $M' \leftarrow \max(M, \max(s))$ ,  $P \leftarrow \exp(S - M')$ , local sum exp  $L1 \leftarrow \sum(P)$ 
        ...

        # Update  $L \leftarrow L * \exp(M - M') + L1$ ,  $M \leftarrow M'$ 
        ...

        # Update  $O \leftarrow O * \exp(M - M') + P @ V$ 
        ...

    return o, m, l

```

# Reading Materials

- [Introduction to torch.compile — PyTorch Tutorials 2.8.0+cu128 documentation](#)
- [Layer Normalization — Triton documentation](#)
- [\[2407.02490\] MInference 1.0: Accelerating Pre-filling for Long-Context LLMs via Dynamic Sparse Attention](#)